

PLSC 30600 — Homework 1

Due: Fri, January 16 (11:59pm)

General instructions

- **Show your work** for all hand calculations. A final numeric answer without intermediate steps, formulas, and/or justification is not sufficient.
 - For coding problems, submit a reproducible script (.R/.rmd/.rnw) along with the fully compiled pdf. Set a random seed and report it.
 - Include a brief note describing how you used AI tools (if at all), consistent with the course AI policy.
 - For each proof question, start with one sentence naming the technique you are using (e.g., direct proof, Law of Iterated Expectations, bounding, counterexample, construction, contradiction).
-

Problem 0: Proof techniques mini-lab

For each part below: (i) state the technique you are using (e.g., bounding, counterexample, construction, Law of Iterated Expectations), and (ii) write a clean argument. You may use the Law of Iterated Expectations without re-proving it. Use Aronow–Miller proofs as templates for identification arguments. Use Hammack or Velleman for general proof strategies (counterexample, construction, bounding, contrapositive).

- 0.a Bounding argument.** Let \tilde{Y}_i be any completion of the missing outcomes such that $\tilde{Y}_i = Y_i$ when $R_i = 1$ and $\tilde{Y}_i \in [0, 1]$ when $R_i = 0$. Let $\tilde{Y} = \frac{1}{n} \sum_{i=1}^n \tilde{Y}_i$. Prove that

$$\underbrace{\frac{1}{n} \sum_{i=1}^n R_i Y_i}_{\text{impute missing as 0}} \leq \tilde{Y} \leq \underbrace{\frac{1}{n} \sum_{i=1}^n (R_i Y_i + (1 - R_i))}_{\text{impute missing as 1}}.$$

(Interpret these two endpoints as the plug-in Manski bounds.)

- 0.b Law of Iterated Expectations proof (MAR identification).** Assume MAR given X :

$$Y \perp R \mid X, \quad \Pr[R = 1 \mid X = x] > 0 \text{ for each } x \in \{A, B\}.$$

Prove that

$$\mathbb{E}[Y] = \sum_{x \in \{A, B\}} \mathbb{E}[Y^* \mid R = 1, X = x] \Pr[X = x].$$

(Your proof should use the Law of Iterated Expectations and the MAR assumption at the appropriate step.)

- 0.c Counterexample (independence vs conditional independence).** Consider binary random variables $X, Y, R \in \{0, 1\}$ with $\Pr[X = 0] = \Pr[X = 1] = 1/2$ and

$$f_{Y,R|X}(y, r \mid x = 0) = \begin{cases} 1/2 & \text{if } (y, r) = (1, 1), \\ 1/2 & \text{if } (y, r) = (0, 0), \\ 0 & \text{otherwise,} \end{cases} \quad f_{Y,R|X}(y, r \mid x = 1) = \begin{cases} 1/2 & \text{if } (y, r) = (1, 0), \\ 1/2 & \text{if } (y, r) = (0, 1), \\ 0 & \text{otherwise.} \end{cases}$$

- (i) Show that $Y \perp R$ (unconditional independence).
- (ii) Show that $Y \not\perp R | X$ (conditional independence fails).

Conclude: unconditional independence does *not* imply conditional independence.

Setup and toy dataset (used in Problem 1 and Coding A–B)

We observe i.i.d. draws of (Y_i^*, R_i, X_i) for $i = 1, \dots, n$, where:

- $Y_i \in [0, 1]$ is the outcome of interest (bounded).
- $R_i \in \{0, 1\}$ is the **response/observation indicator**: $R_i = 1$ means Y_i is observed, $R_i = 0$ means Y_i is missing.
- Y_i^* is the recorded outcome:

$$Y_i^* = \begin{cases} Y_i, & R_i = 1, \\ -99, & R_i = 0. \end{cases}$$

- $X_i \in \{A, B\}$ is a covariate observed for all units (including nonrespondents).

The toy dataset is:

i	X_i	R_i	Y_i^*
1	A	1	0.20
2	A	0	-99
3	B	1	0.80
4	B	0	-99
5	A	1	0.25
6	A	0	-99
7	B	1	0.85
8	B	1	0.90
9	A	0	-99
10	A	1	0.30

Throughout, let $\mu \equiv E[Y]$ denote the population mean of Y .

Problem 1: Manski bounds, MCAR, and MAR on the same dataset

1.a Manski bounds. Assume only that $Y \in [0, 1]$. Compute the **plug-in** lower and upper bounds for $\mu = E[Y]$ using the estimators from Aronow & Miller Theorem 6.1.4.

1.b MCAR (point identification). Now assume **MCAR**:

$$Y \perp R \quad \text{and} \quad \Pr[R = 1] > 0.$$

Compute the MCAR plug-in estimator $\hat{\mu}_{\text{MCAR}}$ from the toy dataset.

1.c MAR given X (point identification by stratification). Now assume **MAR given X** :

$$Y \perp R | X, \quad \Pr[R = 1 | X = x] > 0 \text{ for each } x \in \{A, B\}.$$

Compute the MAR plug-in estimator $\hat{\mu}_{\text{MAR}}$ from the toy dataset.

1.d Interpretation (short answer).

- (i) Why is MCAR typically regarded as *stronger* than MAR?
- (ii) How does the positivity condition differ in spirit between MCAR and MAR?

1.e Sanity checks via bounding (proof). Using the result from Problem 0(a), prove that both $\hat{\mu}_{\text{MCAR}}$ and $\hat{\mu}_{\text{MAR}}$ (from this dataset) must lie inside the plug-in Manski bounds you computed. (Do *not* argue by plugging in numbers; give an inequality argument.)

Problem 2: Propensity scores for missing data

Let $p_R(X) \equiv \Pr[R = 1 | X]$ and assume MAR given X : $Y \perp R | X$.

2.a Balance (conditioning on the propensity score). Show that

$$R \perp X | p_R(X).$$

Hint: because R is binary, it suffices to show $E[R | X, p_R(X)] = E[R | p_R(X)]$.

2.b Outcome-response independence given the propensity score. Show that

$$Y \perp R | p_R(X).$$

Hint: again use the binary- R trick and the Law of Iterated Expectations; under MAR, $E[R | Y, X] = E[R | X] = p_R(X)$.

Problem 3: MCAR vs. MAR diagnostic thinking (vignette)

A city surveys residents to estimate turnout in the most recent election. Let $Y \in \{0, 1\}$ indicate whether a person voted, $R \in \{0, 1\}$ indicate survey response, and let X be age group (recorded for everyone sampled from administrative data, including nonrespondents).

The director believes:

- older residents are more likely to respond (so R depends on X),
- turnout differs by age (so Y depends on X).

3.a Is MCAR ($Y \perp R$) plausible? Why or why not? (2–4 sentences)

3.b Give a realistic story under which MAR given age ($Y \perp R | X$) might be plausible. What does the story rule out? (3–5 sentences)

3.c Suppose age is *only* observed for respondents (you do not observe X when $R = 0$). Which step in the MAR identification argument fails, and why? (3–5 sentences)

Problem 4: Potential outcomes and the basic identification gap (Lemma 1)

Let $D \in \{0, 1\}$ denote a treatment indicator. Let $Y(1)$ and $Y(0)$ be potential outcomes, and let the observed outcome be

$$Y = Y(1)D + Y(0)(1 - D).$$

Lemma 1 (Identification gap). In general,

$$E[Y(1)] \neq E[Y | D = 1] \quad \text{and} \quad E[Y(0)] \neq E[Y | D = 0].$$

4.a Exercise: Prove Lemma 1. Provide a proof by constructing a counterexample distribution for $(Y(0), Y(1), D)$ where at least one inequality holds. (You may use a finite “science table” counterexample or a probability model.)

4.b When does equality hold? State a sufficient condition under which

$$E[Y(1)] = E[Y | D = 1] \quad \text{and} \quad E[Y(0)] = E[Y | D = 0].$$

Briefly justify your answer (a short argument is sufficient).

4.c Selection bias decomposition. Starting from $E[Y | D = 1] - E[Y | D = 0]$, show that

$$E[Y | D = 1] - E[Y | D = 0] = E[Y(1) - Y(0) | D = 1] + \left(E[Y(0) | D = 1] - E[Y(0) | D = 0] \right).$$

Interpret the two terms in 2–4 sentences.

Problem 5: Continuous treatment and average marginal causal effects

Let D be a continuous treatment with support $[0, 1]$ and suppose each unit has potential outcomes $\{Y(d) : d \in [0, 1]\}$, with $Y(d)$ differentiable in d . Assume strong ignorability:

$$\{Y(d)\}_{d \in [0,1]} \perp D | X, \quad 0 < \Pr[D \in \mathcal{N} | X] \text{ for neighborhoods } \mathcal{N} \subset [0, 1].$$

5.a Define the conditional response function $m(d, x) = E[Y | D = d, X = x]$. Show that under strong ignorability,

$$E[Y(d)] = E[m(d, X)].$$

5.b Suppose

$$m(d, x) = \alpha + \beta d + \gamma d^2 + \delta x + \eta(d \cdot x),$$

where $x \in \{0, 1\}$. Compute $\frac{\partial}{\partial d} m(d, x)$.

5.c The (population) average marginal causal effect (AMCE) is defined as

$$\text{AMCE} \equiv E \left[\frac{\partial}{\partial d} m(D, X) \right].$$

Assume $D \sim \text{Uniform}(0, 1)$ and $\Pr[X = 1] = p$ with $D \perp X$. Compute AMCE in terms of (β, γ, η, p) . (Show your work, including any integrals you use.)

Coding section

Your code should be readable and reproducible. You may use Claude Code for assistance, but you are responsible for correctness and interpretation.

Coding A: Replicate Problem 1 numerically

- A1. Enter the toy dataset into your code (as a data frame with columns i , X , R , Y_{star}).

```
toy <- data.frame(  
  i = 1:10,  
  X = c("A", "A", "B", "B", "A", "B", "B", "A", "A"),  
  R = c(1, 0, 1, 0, 1, 0, 1, 1, 0, 1),  
  Ystar = c(0.20, -99, 0.80, -99, 0.25, -99, 0.85, 0.90, -99, 0.30)  
)
```

- A2. Write a function that takes (Y^*, R) and returns the plug-in Manski bounds for $E[Y]$ when $Y \in [0, 1]$.

- A3. Write a function that computes $\hat{\mu}_{\text{MCAR}}$.

- A4. Write a function that computes $\hat{\mu}_{\text{MAR}}$ by stratification on $X \in \{A, B\}$.

- A5. Print the three results (bounds, MCAR estimate, MAR estimate) and verify they match your hand calculations.

Coding B: Bootstrap confidence intervals

Use the nonparametric bootstrap (resample rows with replacement). Letting B be the number of bootstrap resamples, use $B = 100$.

- B1. For each bootstrap sample, compute:

- the lower and upper Manski bounds;
- $\hat{\mu}_{\text{MCAR}}$;
- $\hat{\mu}_{\text{MAR}}$.

(Do *not* print out all 100, just show your code.)

- B2. **Positivity diagnostics in the bootstrap (MAR).** In some bootstrap resamples, one of the strata (A or B) may have no respondents ($R = 1$), making $\hat{\mu}_{\text{MAR}}$ undefined.

- Report the fraction of bootstrap resamples in which $\hat{\mu}_{\text{MAR}}$ is undefined.
- For the resamples where it is defined, compute a 95% percentile CI for μ under MAR.

- B3. Compute 95% percentile bootstrap CIs for:

$$\text{Lower bound}, \quad \text{Upper bound}, \quad \hat{\mu}_{\text{MCAR}}, \quad \hat{\mu}_{\text{MAR}}.$$

Proof toolbox (how to write the proofs in this homework)

Several questions ask for short proofs. You do *not* need advanced real analysis. Most proofs in this course use a small set of reusable techniques. For each proof question:

- start with a one-line label: **Technique:** direct proof / Law of Iterated Expectations / bounding / counterexample / construction / contradiction,
- then write a short, logically complete argument (a clear sequence of equalities/inequalities or implications).

1) Direct proof (“assume the conditions, derive the claim”)

- **When to use:** the statement is of the form “If assumptions A hold, then conclusion C holds.”
- **Template:** Assume A. Start from the left-hand side of what you want to show. Use definitions and algebra (and known facts like linearity of expectation) to rewrite it step-by-step until you obtain the desired expression. End with \square .

2) Law of Iterated Expectations (LIE)

- **Key fact:** for any random variables Y, X ,

$$E[Y] = E(E[Y | X]).$$

- **When to use:** identification arguments where you want to rewrite an unconditional target (like $E[Y]$ or $E[Y(d)]$) using conditional expectations you can relate to observables.
- **Typical move:** write $E[Y] = E(E[Y | X])$, then use an assumption (e.g., MAR/ignorability) to replace $E[Y | X]$ with something observable like $E[Y^* | R = 1, X]$.

3) Bounding (and “sharpness” via extremal cases)

- **When to use:** the target contains an unknown quantity that is only known to lie in an interval.
- **Template:** Rewrite the target so the unknown piece is explicit (often via LIE), then replace the unknown term by its minimum to get a lower bound and by its maximum to get an upper bound.
- **Sharpness:** after deriving bounds, show they are the best possible by *constructing* two data-generating processes (or two completions of missing outcomes) that attain the lower and upper bounds.

4) Counterexample (disproving a false general claim)

- **When to use:** the claim says “for all” (e.g., “in general $E[Y(1)] = E[Y | D = 1]$ ”).
- **Template:** give one explicit joint distribution (or a small “science table”) where the assumptions hold but the conclusion fails. One example is enough.

5) Construction (showing existence or attaining an extremum)

- **When to use:** you must show something *can* happen (e.g., sharpness, or that two DGPs are observationally equivalent).
- **Template:** explicitly define the object you need (e.g., fill in missing Y values, or define a DGP for $Y | R = 0$), then verify it matches the observed data/assumptions.

6) Proof by contradiction (assume the opposite, derive an impossibility)

- **When to use:** to show a condition is necessary, or to show an event is impossible.
- **Template:** assume the negation of the statement you want, then derive a contradiction (e.g., a probability < 0 or > 1 , or two different values for something that should be unique).

Common “allowed moves” in this course

- Linearity: $E[aY + bW] = aE[Y] + bE[W]$ (and similarly for conditional expectations).
- Conditioning on a function adds no information if you already condition on the variable: $E[U | X, g(X)] = E[U | X]$.
- If $R \in \{0, 1\}$ is binary, its conditional distribution is determined by $E[R | \cdot]$. So to prove a conditional independence involving R , it often suffices to show an equality of conditional expectations.

Where you have already seen these in the reading

- Bounding + sharpness for missing data (Manski bounds): see the sharp bounds theorem and proof.
- Law of Iterated Expectations + substitution under MAR/ignorability: see the MAR identification proofs.
- Binary-variable trick + conditioning arguments: see the propensity score theorem proof.
- Counterexample using a science table: see the post-treatment-variable example.