

PLSC 30600

Week 6: Design-based inference

Molly Offer-Westort

Department of Political Science,
University of Chicago

Winter 2026

Loading packages for this class

```
> # install.packages('../assets/ri_0.9.tar.gz',  
> # repos = NULL, type = 'source')  
> library(ri)  
> library(ggplot2)  
> set.seed(60637)
```

Example adapted from:

Chattopadhyay, Raghavendra & Duflo, Esther (2004). Women as policy makers: Evidence from a randomized policy experiment in India.

as discussed in

Gerber, Alan S. & Green, Donald P. (2012). Field experiments: Design, analysis, and interpretation.

Chattopadhyay & Duflo example

- Population: 7 villages

Chattopadhyay & Duflo example

- Population: 7 villages
- Treatment: $D = 1$ if female-headed council, $D = 0$ if male

Chattopadhyay & Duflo example

- Population: 7 villages
- Treatment: $D = 1$ if female-headed council, $D = 0$ if male
- Outcome: Budget allocation to sanitation

Chattopadhyay & Duflo example

- Population: 7 villages
- Treatment: $D = 1$ if female-headed council, $D = 0$ if male
- Outcome: Budget allocation to sanitation

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	10	15	1	15	5
Village 2	15	15	0	15	0
Village 3	20	30	0	20	10
Village 4	20	15	0	20	-5
Village 5	10	20	0	10	10
Village 6	15	15	0	15	0
Village 7	15	30	1	30	15
Average	15	20			5

$$ATE = E[\tau_i] = 5$$

Chattopadhyay & Duflo example

What we actually see:

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	?	15	1	15	?
Village 2	15	?	0	15	?
Village 3	20	?	0	20	?
Village 4	20	?	0	20	?
Village 5	10	?	0	10	?
Village 6	15	?	0	15	?
Village 7	?	30	1	30	?

Chattopadhyay & Duflo example

What we actually see:

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	?	15	1	15	?
Village 2	15	?	0	15	?
Village 3	20	?	0	20	?
Village 4	20	?	0	20	?
Village 5	10	?	0	10	?
Village 6	15	?	0	15	?
Village 7	?	30	1	30	?

To produce an *estimate* of the ATE, we can compare people who received treatment 1 to people who received treatment 0.

Chattopadhyay & Duflo example

What we actually see:

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	?	15	1	15	?
Village 2	15	?	0	15	?
Village 3	20	?	0	20	?
Village 4	20	?	0	20	?
Village 5	10	?	0	10	?
Village 6	15	?	0	15	?
Village 7	?	30	1	30	?

To produce an *estimate* of the ATE, we can compare people who received treatment 1 to people who received treatment 0.

$$\bar{Y}(1) = 22.5 \quad \bar{Y}(0) = 16$$

Chattopadhyay & Duflo example

What we actually see:

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	?	15	1	15	?
Village 2	15	?	0	15	?
Village 3	20	?	0	20	?
Village 4	20	?	0	20	?
Village 5	10	?	0	10	?
Village 6	15	?	0	15	?
Village 7	?	30	1	30	?

To produce an *estimate* of the ATE, we can compare people who received treatment 1 to people who received treatment 0.

$$\bar{Y}(1) = 22.5 \quad \bar{Y}(0) = 16$$

$$\hat{\tau} = 22.5 - 16 = \mathbf{6.5}$$



estimand



estimate

Ingredients

150g unsalted butter, plus extra for greasing
150g plain chocolate, broken into pieces
150g plain flour
¼ tsp baking powder
¼ tsp bicarbonate of soda
200g light muscovado sugar
2 large eggs

Method

1. Heat the oven to 160C/140C fan/gas 3. Grease and base line a 1 litre heatproof glass pudding basin and a 450g loaf tin with baking parchment.
2. Put the butter and chocolate into a saucepan and melt over a low heat, stirring. When the chocolate has all melted remove from the heat.

estimator

(image cred: simongrund89)

Difference-in-means

Proposed estimator for $E[\tau_i]$: compare people who received treatment 1 to people who received treatment 0.

$$\hat{\tau}_{DM} = \frac{\sum_i^n Y_i D_i}{\sum_i^n D_i} - \frac{\sum_i^n Y_i (1 - D_i)}{\sum_i^n (1 - D_i)}$$

Chattopadhyay & Duflo example

- We have produced an estimate here, by taking the difference in means.

Chattopadhyay & Duflo example

- We have produced an estimate here, by taking the difference in means.
- But what can we say about our uncertainty about this estimate? The number of observations in each group is pretty small.

Chattopadhyay & Duflo example

- We have produced an estimate here, by taking the difference in means.
- But what can we say about our uncertainty about this estimate? The number of observations in each group is pretty small.
- Is the estimate meaningfully different from zero?

Chattopadhyay & Duflo example

- We have produced an estimate here, by taking the difference in means.
- But what can we say about our uncertainty about this estimate? The number of observations in each group is pretty small.
- Is the estimate meaningfully different from zero?
- How likely would we be to see an effect this size just by chance?

Chattopadhyay & Duflo example

- One way to think about this is to assume the individual treatment effect for all individuals is exactly zero. Then, no matter how we randomized treatment assignment, we would see the same Y_i s.

Chattopadhyay & Duflo example

- One way to think about this is to assume the individual treatment effect for all individuals is exactly zero. Then, no matter how we randomized treatment assignment, we would see the same Y_i s.
- We can then say how often we would see a treatment effect estimate of this size just by chance, under the assumption that individual treatment effects were actually zero.

Sharp null hypothesis of no effect

Assuming all treatment effects are exactly zero is called the **sharp null hypothesis of no effect**.

Sharp null hypothesis of no effect

Assuming all treatment effects are exactly zero is called the **sharp null hypothesis of no effect**. Also referred to as “Fisher’s null” after Sir Ronald Aylmer Fisher (1890-1962).



Sharp null hypothesis of no effect

We might write the sharp null hypothesis this way:

$$H_0 : \tau_i = 0, \text{ for all } i \text{ in our pop.}$$

This implies that potential outcomes are identical under treatment and control, for all individuals.

$$Y_i(0) = Y_i(1)$$

Sharp null hypothesis of no effect

- To test this hypothesis, we also need to know (or assume we know) the randomization procedure

Sharp null hypothesis of no effect

- To test this hypothesis, we also need to know (or assume we know) the randomization procedure
- Here, we'll assume exactly two villages are assigned treatment.

Chattopadhyay & Duflo example

Then we can fill in this table...

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	?	15	1	15	?
Village 2	15	?	0	15	?
Village 3	20	?	0	20	?
Village 4	20	?	0	20	?
Village 5	10	?	0	10	?
Village 6	15	?	0	15	?
Village 7	?	30	1	30	?

Chattopadhyay & Duflo example

...as this:

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	1	15	0
Village 2	15	15	0	15	0
Village 3	20	20	0	20	0
Village 4	20	20	0	20	0
Village 5	10	10	0	10	0
Village 6	15	15	0	15	0
Village 7	30	30	1	30	0

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	1	15	0
Village 3	20	20	0	20	0
Village 4	20	20	0	20	0
Village 5	10	10	0	10	0
Village 6	15	15	1	15	0
Village 7	30	30	0	30	0

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	1	15	0
Village 3	20	20	0	20	0
Village 4	20	20	0	20	0
Village 5	10	10	0	10	0
Village 6	15	15	1	15	0
Village 7	30	30	0	30	0

$$\bar{Y}(1) = 15 \quad \bar{Y}(0) = 19$$

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	1	15	0
Village 3	20	20	0	20	0
Village 4	20	20	0	20	0
Village 5	10	10	0	10	0
Village 6	15	15	1	15	0
Village 7	30	30	0	30	0

$$\bar{Y}(1) = 15 \quad \bar{Y}(0) = 19$$

$$\hat{\tau} = 15 - 19 = -4$$

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	0	15	0
Village 3	20	20	1	20	0
Village 4	20	20	1	20	0
Village 5	10	10	0	10	0
Village 6	15	15	0	15	0
Village 7	30	30	0	30	0

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	0	15	0
Village 3	20	20	1	20	0
Village 4	20	20	1	20	0
Village 5	10	10	0	10	0
Village 6	15	15	0	15	0
Village 7	30	30	0	30	0

$$\bar{Y}(1) = 20 \quad \bar{Y}(0) = 17$$

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	0	15	0
Village 3	20	20	1	20	0
Village 4	20	20	1	20	0
Village 5	10	10	0	10	0
Village 6	15	15	0	15	0
Village 7	30	30	0	30	0

$$\bar{Y}(1) = 20 \quad \bar{Y}(0) = 17$$

$$\hat{\tau} = 20 - 17 = 3$$

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	0	15	0
Village 3	20	20	1	20	0
Village 4	20	20	0	20	0
Village 5	10	10	0	10	0
Village 6	15	15	0	15	0
Village 7	30	30	1	30	0

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	0	15	0
Village 3	20	20	1	20	0
Village 4	20	20	0	20	0
Village 5	10	10	0	10	0
Village 6	15	15	0	15	0
Village 7	30	30	1	30	0

$$\bar{Y}(1) = 25 \quad \bar{Y}(0) = 15$$

Gerber Green example

We can re-run the randomization, and the potential outcomes and observed Y_i will not change, but the treatment effect estimate will.

	$Y_i(0)$	$Y_i(1)$	D_i	Y_i	τ_i
Village 1	15	15	0	15	0
Village 2	15	15	0	15	0
Village 3	20	20	1	20	0
Village 4	20	20	0	20	0
Village 5	10	10	0	10	0
Village 6	15	15	0	15	0
Village 7	30	30	1	30	0

$$\bar{Y}(1) = 25 \quad \bar{Y}(0) = 15$$

$$\hat{\tau} = 25 - 15 = 10$$

Randomization procedure

- Design: complete randomization with $N = 7$ villages and exactly $n_1 = 2$ treated.

Randomization procedure

- Design: complete randomization with $N = 7$ villages and exactly $n_1 = 2$ treated.
- Each assignment vector $\mathbf{D}$ with $\sum_i D_i = 2$ is equally likely:

$$\Pr[\mathbf{D} = \mathbf{d}] = \frac{1}{\binom{7}{2}} = \frac{1}{21}.$$

Randomization procedure

- Design: complete randomization with $N = 7$ villages and exactly $n_1 = 2$ treated.
- Each assignment vector $\mathbf{D}$ with $\sum_i D_i = 2$ is equally likely:

$$\Pr[\mathbf{D} = \mathbf{d}] = \frac{1}{\binom{7}{2}} = \frac{1}{21}.$$

- The marginal assignment probability is the same for each village:

$$\pi_i \equiv \Pr[D_i = 1] = \frac{2}{7}.$$

We can use the package `ri` to find the *exact* distribution of $\hat{\tau}_{DM}$ under the sharp null.

Our real data:

```
> df <- data.frame(  
+ # our initial treatment vector  
+ D = c(1, 0, 0, 0, 0, 0, 1),  
+ # our initial response vector  
+ Y = c(15, 15, 20, 20, 10, 15, 30),  
+ # marginal assignment probability:  $\pi_i = \Pr(D_i = 1)$   
+ pi = rep(2/7, 7)  
+ )  
> df
```

	D	Y	pi
1	1	15	0.2857143
2	0	15	0.2857143
3	0	20	0.2857143
4	0	20	0.2857143
5	0	10	0.2857143
6	0	15	0.2857143

And our difference in means estimate of the average treatment effect under the real data:

```
> Y1 ← df$Y[which(df$D == 1)]  
> Y0 ← df$Y[which(df$D == 0)]  
> (dm_hat ← mean(Y1) - mean(Y0))
```

```
[1] 6.5
```

Adding in the hypothetical data.

```
> df ← cbind( # binds the columns together
+ df,
+ # Y(0) under the sharp null of no effect
+ Y0 = df$Y,
+ # Y(1) under the sharp null of no effect
+ Y1 = df$Y)
> df
```

	D	Y	pi	Y0	Y1
1	1	15	0.2857143	15	15
2	0	15	0.2857143	15	15
3	0	20	0.2857143	20	20
4	0	20	0.2857143	20	20
5	0	10	0.2857143	10	10
6	0	15	0.2857143	15	15
7	1	30	0.2857143	30	30

Consider all the ways treatment could be assigned: (this is what `ri` is doing for us)

```
> (perms ← genperms(df$D))
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
1	1	1	1	1	1	1	0	0	0	0	
0	0		0	0							
2	1	0	0	0	0	0	1	1	1	1	
1	0		0	0							
3	0	1	0	0	0	0	1	0	0	0	
0	1		1	1							
4	0	0	1	0	0	0	0	1	0	0	
0	1		0	0							
5	0	0	0	1	0	0	0	0	1	0	
0	0		1	0							
6	0	0	0	0	1	0	0	0	0	1	
0	0		0	1							
7	0	0	0	0	0	1	0	0	0	0	
1	0		0	0							
	[,15]	[,16]	[,17]	[,18]	[,19]	[,20]	[,21]				
1	0	0	0	0	0	0	0				
2	0	0	0	0	0	0	0				

Then generate the sampling distribution of the ATE estimate under the sharp null of no effect.

```
> Ys_null <- list(  
+ Y0 = df$Y0,  
+ Y1 = df$Y1  
+ )  
> dm <- gendist(Ys_null,  
+ perms,  
+ prob=df$pi)  
> dm
```


```
[1] -4.0 -0.5 -0.5 -7.5 -4.0  6.5 -0.5 -0.5 -7.5 -4.0  
6.5  3.0 -4.0 -0.5 10.0  
[16] -4.0 -0.5 10.0 -7.5  3.0  6.5
```

The mean under our null distribution is *exactly* zero. Why?

```
> mean(dm)
```

```
[1] 0
```

Distribution of difference in means estimates under the sharp null



```
> prop.table(table(dm))
```

```
dm
   -7.5    -4   -0.5    3    6.5
10
0.1428571 0.2380952 0.2857143 0.0952381 0.1428571 0.0952381
```

```
> (pval <- mean(abs(dm) ≥ abs(dm_hat)))
```

```
[1] 0.3809524
```

p-value

- Under the null distribution, if we were to re-randomize the experiment many times, we would see a value *at least as extreme as our estimate* 38.1% of the time.

p-value

- Under the null distribution, if we were to re-randomize the experiment many times, we would see a value *at least as extreme as our estimate* 38.1% of the time.
- That doesn't seem **that** unlikely—more than one in three times, we would see an estimate as large as we got, just by chance.

ri produces the same exact p-value.

```
> dispdist(distout = dm,  
+ ate = dm_hat,  
+ display.plot = FALSE)$two.tailed.p.value.abs
```

```
[1] 0.3809524
```

ri produces the same exact p-value.

```
> dispdist(distout = dm,  
+ ate = dm_hat,  
+ display.plot = FALSE)$two.tailed.p.value.abs
```

```
[1] 0.3809524
```

That's because it is doing the exact same thing under the hood.

- Why is this test called an “exact” test?

- Why is this test called an “exact” test?
- Because we know the *exact* distribution of our estimate under the specified null. We do not have to approximate the distribution.
- This will not be true for all of our hypothesis testing...

Hypotheses

- We framed our null hypothesis as below:

$$H_0 : \tau_i = 0, \text{ for all } i \text{ in our pop.}$$

Hypotheses

- We framed our null hypothesis as below:

$$H_0 : \tau_i = 0, \text{ for all } i \text{ in our pop.}$$

- Implicitly, the alternative is that for some individual(s), the treatment effect is non-zero.

$$H_A : \tau_i \neq 0, \text{ for some } i \text{ in our pop.}$$

Hypotheses

- We framed our null hypothesis as below:

$$H_0 : \tau_i = 0, \text{ for all } i \text{ in our pop.}$$

- Implicitly, the alternative is that for some individual(s), the treatment effect is non-zero.

$$H_A : \tau_i \neq 0, \text{ for some } i \text{ in our pop.}$$

- In our case, we did not find strong evidence to reject the null hypothesis, i.e., our data is consistent with what we would see if the null hypothesis were true.

Hypotheses

- We framed our null hypothesis as below:

$$H_0 : \tau_i = 0, \text{ for all } i \text{ in our pop.}$$

- Implicitly, the alternative is that for some individual(s), the treatment effect is non-zero.

$$H_A : \tau_i \neq 0, \text{ for some } i \text{ in our pop.}$$

- In our case, we did not find strong evidence to reject the null hypothesis, i.e., our data is consistent with what we would see if the null hypothesis were true.
- Note that we do NOT say that we reject or accept the alternative hypothesis.

Hypotheses

- We framed our null hypothesis as below:

$$H_0 : \tau_i = 0, \text{ for all } i \text{ in our pop.}$$

- Implicitly, the alternative is that for some individual(s), the treatment effect is non-zero.

$$H_A : \tau_i \neq 0, \text{ for some } i \text{ in our pop.}$$

- In our case, we did not find strong evidence to reject the null hypothesis, i.e., our data is consistent with what we would see if the null hypothesis were true.
- Note that we do NOT say that we reject or accept the alternative hypothesis.
- We can only say that our results were not consistent with or were consistent with what we would have seen under the null—i.e., we have evidence to reject or fail to reject the null.

Distributions of estimators

- Our difference in means estimator is a function of the data we observe.

Distributions of estimators

- Our difference in means estimator is a function of the data we observe.
- Because there is randomness in the data, here, due to random assignment of treatment, the estimator is also a random variable.

Distributions of estimators

- Our difference in means estimator is a function of the data we observe.
- Because there is randomness in the data, here, due to random assignment of treatment, the estimator is also a random variable.
- Just like other random variables have distributions to describe them, estimators also have distributions.

Distributions of estimators

- Our difference in means estimator is a function of the data we observe.
- Because there is randomness in the data, here, due to random assignment of treatment, the estimator is also a random variable.
- Just like other random variables have distributions to describe them, estimators also have distributions.
- We don't know the *true* distribution of the difference in means estimator, for the same reason that we don't know individual treatment effects. (FPoCI!)

Distributions of estimators

- Our difference in means estimator is a function of the data we observe.
- Because there is randomness in the data, here, due to random assignment of treatment, the estimator is also a random variable.
- Just like other random variables have distributions to describe them, estimators also have distributions.
- We don't know the *true* distribution of the difference in means estimator, for the same reason that we don't know individual treatment effects. (FPoCI!)
- But we DO know what the distribution would be under the null.

Note that we are conducting inference with respect to:

Note that we are conducting inference with respect to:

- a defined population

Note that we are conducting inference with respect to:

- a defined population
- a defined treatment

Note that we are conducting inference with respect to:

- a defined population
- a defined treatment
- a defined outcome

Note that we are conducting inference with respect to:

- a defined population
- a defined treatment
- a defined outcome
- a known treatment assignment mechanism

Note that we are conducting inference with respect to:

- a defined population
- a defined treatment
- a defined outcome
- a known treatment assignment mechanism
- a given estimating procedure

The null

- How do we determine what the null is?

The null

- How do we determine what the null is?
- We formalize our hypotheses in terms of the effect we are trying to find in the data. Is there a treatment effect? Is there a difference between these two groups?

The null

- How do we determine what the null is?
- We formalize our hypotheses in terms of the effect we are trying to find in the data. Is there a treatment effect? Is there a difference between these two groups?
- The null is (often, but not always) the case when there is no effect, or no difference.

The null

- How do we determine what the null is?
- We formalize our hypotheses in terms of the effect we are trying to find in the data. Is there a treatment effect? Is there a difference between these two groups?
- The null is (often, but not always) the case when there is no effect, or no difference.
- We can imagine other kinds of hypotheses, for example that effects are bounded away from zero and positive, or exactly .2. And we can characterize the distribution of our test statistic under the null.

The null

- Is the sharp null of no individual level effects plausible in this setting?

The null

- Is the sharp null of no individual level effects plausible in this setting?
- Does it matter?

Worked example: Fisher's "Lady tasting tea"

Fisher (1935), Chapter 2

Randomization inference by counting assignments

Setup

- 8 cups of tea

Setup

- 8 cups of tea
- 4 cups prepared **milk-first** and 4 prepared **tea-first**

Setup

- 8 cups of tea
- 4 cups prepared **milk-first** and 4 prepared **tea-first**
- The experimenter **randomizes** which cups are which

Setup

- 8 cups of tea
- 4 cups prepared **milk-first** and 4 prepared **tea-first**
- The experimenter **randomizes** which cups are which
- The lady claims she can identify which cups were prepared milk-first

Setup

- 8 cups of tea
- 4 cups prepared **milk-first** and 4 prepared **tea-first**
- The experimenter **randomizes** which cups are which
- The lady claims she can identify which cups were prepared milk-first
- She must label **exactly 4** cups as milk-first (so guessing has a known baseline)

Null + test statistic

- Null (no skill): her answers are unrelated to the randomized assignment

Null + test statistic

- Null (no skill): her answers are unrelated to the randomized assignment
- Encode the randomization as $\mathbf{Z} \in \{0, 1\}^8$, where $Z_i = 1$ if cup i is milk-first and $\sum_i Z_i = 4$

Null + test statistic

- Null (no skill): her answers are unrelated to the randomized assignment
- Encode the randomization as $\mathbf{Z} \in \{0, 1\}^8$, where $Z_i = 1$ if cup i is milk-first and $\sum_i Z_i = 4$
- Encode her guess as $\mathbf{G} \in \{0, 1\}^8$, where $G_i = 1$ if she says cup i is milk-first and $\sum_i G_i = 4$

Null + test statistic

- Null (no skill): her answers are unrelated to the randomized assignment
- Encode the randomization as $\mathbf{Z} \in \{0, 1\}^8$, where $Z_i = 1$ if cup i is milk-first and $\sum_i Z_i = 4$
- Encode her guess as $\mathbf{G} \in \{0, 1\}^8$, where $G_i = 1$ if she says cup i is milk-first and $\sum_i G_i = 4$
- Test statistic (a dot product / contrast):

$$T = \mathbf{Z}^\top \mathbf{G} = \sum_{i=1}^8 Z_i G_i$$

which equals the number of milk-first cups she correctly identifies.

Step 1: Observed data (one realization)

```
> tea <- data.frame(  
+ cup = 1:8,  
+ # randomized assignment: which cups are  
+   milk-first?  
+ Z = c(1, 1, 0, 0, 1, 1, 0, 0),  
+ # her guess: which cups does she label  
+   milk-first?  
+ G = c(1, 1, 1, 0, 1, 0, 0, 0)  
+ )  
> tea$correct <- tea$Z * tea$G  
> tea
```

	cup	Z	G	correct
1	1	1	1	1
2	2	1	1	1
3	3	0	1	0
4	4	0	0	0
5	5	1	1	1
6	6	1	0	0
7	7	0	0	0

Step 2: How many assignments are possible?

Under the design, we choose which 4 of the 8 cups are milk-first:

$$\binom{8}{4} = 70$$

Step 2: How many assignments are possible?

Under the design, we choose which 4 of the 8 cups are milk-first:

$$\binom{8}{4} = 70$$

Each assignment is equally likely under the randomization.

Or generate permutations

```
> (perms ← genperms(tea$Z))
```

[illegible]

Step 3: Exact randomization distribution under the null


```
> dist_df <- data.frame(  
+ T = 0:4,  
+ prob = dhyper(0:4, m = 4, n = 4, k = 4)  
+ )  
> dist_df
```

	T	prob
1	0	0.01428571
2	1	0.22857143
3	2	0.51428571
4	3	0.22857143
5	4	0.01428571

```
> (p_ge_Tobs <- sum(dist_df$prob[dist_df$T ≥ T_obs  
  ]))
```

```
[1] 0.2428571
```

Lady tasting tea: exact randomization distribution under the sharp null



Step 5: Compute and interpret the p-value

In this example, $T_{\text{obs}} = 3$. The one-sided p-value is:

$$p = \Pr[T \geq T_{\text{obs}}] = 0.2429.$$

Step 5: Compute and interpret the p-value

In this example, $T_{\text{obs}} = 3$. The one-sided p-value is:

$$p = \Pr[T \geq T_{\text{obs}}] = 0.2429.$$

- This p-value is computed by **counting assignments** consistent with results at least as good as what we saw.

Step 5: Compute and interpret the p-value

In this example, $T_{\text{obs}} = 3$. The one-sided p-value is:

$$p = \Pr[T \geq T_{\text{obs}}] = 0.2429.$$

- This p-value is computed by **counting assignments** consistent with results at least as good as what we saw.
- It is **not** $P(H_0 \mid \text{data})$ (Fisher rejects “inverse probability” as the general justification for scientific conclusions).

References I

- Chattopadhyay, R. and Duflo, E. (2004). Women as policy makers: Evidence from a randomized policy experiment in india. *Econometrica*, 72(5):1409–1443.
- Fisher, R. A. (1935). *The Design of Experiments*. Oliver and Boyd Ltd., Edinburgh, 1 edition.
- Gerber, A. S. and Green, D. P. (2012). *Field Experiments: Design, Analysis, and Interpretation*. W. W. Norton.